

# FPGA notes

Fabien Le Mentec  
fabien.lementec@gmail.com

version: 4a736f6

## **Abstract**

Set of FPGA related notes, written by a developer still learning about FPGAs in general, so any feedback welcome. It addresses topics ranging from coding conventions, verification, synthesis, optimisation, reusability and documentation. Some notes are vague while others are quite specific to language, tools or target platforms. Also, it addresses a wide audience, so some materials may seem obvious to the reader, depending on her background.

# Contents

1	Unconstrained types	3
2	Type attributes	5
3	Generics instead of package constants	6
4	Per component test benches	8
5	Hardware resource inference	9
6	Explicit resource instantiation	10
7	Reset signals	11
8	Shift registers inference	12
9	Assertions	13
10	Test benches as documentation	14
11	Writing synchronous processes	15
12	Clocking	16
13	Appropriate typing	17
14	Component directory structure	18
15	Use OR for multiplexers	19
16	Variable length multiplexers	20
17	Modular simulation	21
18	Include physics in simulations	23

# 1 Unconstrained types

*topics: reusability, documentation, verification*

When not explicitly specified by the developer, a type length is deduced during component instantiation. This favors component reusability by letting the user decide of the type length according to its particular needs.

For instance, *count* range is unconstrained in the following component declaration:

```
component my_component
port
(
  ...
  count: in unsigned;
  ...
);
```

The actual range is deduced during instantiation:

```
signal count: unsigned(7 downto 0);

work.my_component
port map
(
  ...
  count => count
  ...
);
```

One important issue with unconstrained types is that a component user may inadvertently use types that are larger than required, possibly leading to unnecessary large resource instantiation. Documentation is a good tool to solve this kind of issue. Also, assertions can be used to check for degenerate cases:

```

component my_component
port
(
    ...
    — WARNING
    — hardware comparator inferred in subsequent logic
    — use appropriate length
    count: in unsigned;
    ...
);
end component;

```

If you want to make sure the length fits within a given range, use assertion in the component entity definition:

```

entity my_component
port
(
    ...
    count: in unsigned;
    ...
);
end my_component;

architecture my_component_rtl of my_component is
begin
    ...
    assert (count'length <= 16)
    report "invalid_counter_length"
    severity failure;
    ...
end my_component_rtl;

```

## 2 Type attributes

Use type attribute as much as possible, esp. type'length and type'range

### 3 Generics instead of package constants

*topics: reusability*

Often, a component parameter can be set either using package constants or generics. Using package constants forces the user to modify the package source code, which is a bad thing considering a package should be reusable without any modification. On the other hand, generics let the user specialize the component without modifying any existing source.

In the following example, the clock frequency (CLK\_FREQ) is used by the component called slave. The frequency could be defined as a constant in the package, like this:

```
package abs_enc_pkg is
...
constant CLK_FREQ: integer := 12500000;
...
component slave
port
(
  clk: in std_logic;
  rst: in std_logic;
  ...
);
end component;
...
end package abs_enc_pkg;
```

A user wanting to change CLK\_FREQ the clock frequency must change the abs\_enc\_pkg source file. This can be avoided by defining CLK\_FREQ as a component generic:

```
package abs_enc_pkg is  
component slave  
generic  
(  
  CLK_FREQ: integer := 12500000;  
);  
port  
(  
  clk: in std_logic;  
  rst: in std_logic;  
  ...  
);  
end component;  
...  
end package abs_enc_pkg;
```

Since the CLK\_FREQ is not globally visible, it must be propagated through the design hierarchy. This might be inconvenient as the generics lists grow larger, but a clear advantage for reusability.

**TODO:** Global variables may be used to solve this issue, but I do not think their use is encouraged.

## 4 Per component test benches

*topics: simulation*

Per component test benches generally requires less code than project wide ones. It makes them easier to maintain, and encourages the developer to write self contained components. Also, it makes simulation run faster.



## 5 Hardware resource inference

*topics: synthesis*

Usually, a VHDL developer does not explicitly indicate what hardware resource to use to implement logic. The synthesiser deduces that from its source code understanding (ie. signal netlist and operations). This process is known as inference.

Inference is very sensitive to the way code is written. For instance, the use of an additional signal to reset a shift register may prevent the synthesiser to infer a hardware shift register.

Thus, VHDL developers try as much as possible to write code in a standard way, that is known to be well understood by the synthesiser.

## 6 Explicit resource instantiation

*topics: synthesis*

**TODO:** wip

Non portable but sure to instanciate the right resource.

## 7 Reset signals

*related notes: 11*

Avoid reset signals. If not possible, make reset synchronous.

**TODO:** explain why

## 8 Shift registers inference

*topics: synthesis*

**TODO:** wip

XILINX FPGAs have hardware resources to implement shift registers. Reference here (p.78):  
[http://www.xilinx.com/support/documentation/sw\\_manuals/xilinx2014\\_1/ug901-vivado-synthesis.pdf](http://www.xilinx.com/support/documentation/sw_manuals/xilinx2014_1/ug901-vivado-synthesis.pdf)

## 9 Assertions

*topics: verification*

**TODO:** wip

Use assertion to check data type lengths when unconstraints arrays

## 10 Test benches as documentation

*topics: documentation*

A component developer should consider test benches an important part of the documentation since they are used as reference materials by the component user. Thus, test benches should be up to date, clearly written and well documented. If possible, they should cover different use cases, without flooding the user with unrequired contents.

## 11 Writing synchronous processes

*topics: synthesis*

There is one standard way of writing synchronous process:

```
process (clk , rst)
begin
  if rising_edge (clk) then
    if rst = '1' then
      else
    end if;
  end if;
end process;
```

Another way which is synthetizable:

```
process
begin
  wait until rising_edge (clk);

  if rst = '1' then
    end if;

end process;
```

Since the **wait** statement must come first, all the signal are synchronous, esp. the reset. Also, this convention results in a somewhat clearer code.

## 12 Clocking

**TODO:** wip

Clear convention about how data passed to/from a component are clocked. by default, clocked using the component domain. idem for latching.



## 13 Appropriate typing

*topics: verification, documentation*

Use the most specialized types (unsigned, boolean ...) and sizes early in the design hierarchy. It avoids further casting and simplifies the code. It acts as documentation since the reader deduces information from the type itself. For instance, an unsigned counter tells it can not be negative. Typing also improve static time checks.

## 14 Component directory structure

*topics: reusability, documentation*

Having a clear, self contained directory structure is helpful for both the user and the developer. I opted for the following one, that simple but fits most of the cases:

```
my_component /
src /
  my_component_pkg.vhd
  my_component_rtl.vhd
  ...
sim /
  common /
    main_tb.vhd
  isim /
    isim.tcl
    isim.prj
    isim.sh
  modelsim /
    ...
syn /
  ise /
    xc7k325t.ucf
    xc7vx485t.ucf
  vivado /
    ...
doc /
  my_component.pdf
```

Providing synthesis files allows the user to synthesise the component for a given platform. It should not synthesise a full working design, only the bare minimum so the user can check its toolchain (esp. version), and investigate what hardware resources are inferred.

## 15 Use OR for multiplexers

*topics: optimization*

**TODO:** check if it really optimize, explain

## 16 Variable length multiplexers

*topics: reusability, optimization*

**TODO:**

## 17 Modular simulation

*topics: reusability, simulation*

**TODO:** add composed simulation scenarios

Generally, a simulation is implemented as a single monolithic file. Signals are commented based on what the user wants to investigate. This approach is inconvenient at least, especially when signals are mutually exclusive.

Another approach is to split the simulation into multiple scripts. A modular environment favors new scenario composition by aggregating existing scripts. Also, it makes it easier for a user to investigate only the signals he is interested in. Also, it makes the environment easier for the developer to maintain and reuse.

For instance, a project implementing a generic serial link controller may use the following structure for an ISIM simulation:

```
generic_serial_controller/  
...  
sim/  
  common/  
    main.vhd  
  isim/  
    i2c.tcl  
    spi.tcl  
    onewire.tcl  
    common.tcl  
    main.tcl  
...
```

In this example, the main file is in charge of sourcing other files and running the simulation:

```

source common.tcl

if { [ file exists user.tcl ] == 1 } {
    source user.tcl
} else {
    source i2c.tcl
    source spi.tcl
    source onewire.tcl
}

run 150 us

```

The file common.tcl adds the signals that will be required in any simulation. What signal is required is left as a decision to the developer. In this example, both system clock, reset and serial link clock and data are added:

```

isim force add {/main/clock} \
1 -value 0 -radix bin -time 10 ns -repeat 20 ns

isim force add {/main/rst} \
1 -value 0 -time 2 us

wave add /main/master_clock
wave add /main/master_data
wave add /main/slave_data

```

The other files add protocol specific signal. For instance, spi.tcl:

```

wave add /main/master/gen_spi/master_spi/curr_state
wave add /main/master/gen_spi/master_spi/spi_clk
wave add /main/master/gen_spi/master_spi/spi_miso
wave add /main/master/gen_spi/master_spi/spi_mosi

```

Finally, a user specific file is included if it exists. It may contain whatever signal is of interest to the user. With this approach, there is no need for the user to modify existing files.

## 18 Include physics in simulations

*topics: simulation*

**TODO:** signal propagation delays due to cable length

**TODO:** noise sources, debouncers